

## GPGPU

// Variables

\_\_device\_\_

- Stored in global memory (large, high latency, no cache)
- Read/write by all threads within grid
- Written by CPU via cudaMemcpyToSymbol()
- Lifetime: application

\_\_constant\_\_

- Same as \_\_device\_\_, but cached and read-only by all threads within grid
- Written by CPU via cudaMemcpyToSymbol()
- Lifetime: application

\_\_shared\_\_

- Stored in on-chip shared memory (very low latency)
- Read/write by all threads in the same thread block
- Lifetime: block

Unqualified variables (in device code)

- Scalars and built-in vector types are stored in registers
- Arrays of more than 4 elements or run-time indices stored in local memory
- Read/write by thread only
- Lifetime: thread

// must be volatile to use array as shared memory value  
otherwise might get optimized away into registers  
\_\_device\_\_ \_\_managed\_\_ volatile int global\_counter[2];

// global memory on device

\_\_device\_\_ \_\_managed\_\_ int counter;

\_\_global\_\_ func can be called by host or device

\_\_device\_\_ can only be called by device

no recursion

no return value --> data transfer by cudaMemcpy

printf can come from global func!

// atomic operations

atomicAdd(&counter, 1);

// calling Syntax

non\_atomic<<<gridDim, blockDim>>>();

kernel<<<dim3 grid, dim3 block, int smem, int stream>>>(...)

mykernel<<<blocks, threads, shared\_mem, stream>>>(args);

// synchronize different streams

in host --> cudaDeviceSynchronize();

barrier in block --> \_\_syncthreads();

// memory --> copy the data from address start to address device\_mem

// allocate memory on device

cudaMalloc((void \*\*)&device\_mem, num\_elements \* sizeof(int));

// host to device

// target, source, size

rc = cudaMemcpy(device\_mem, start, num\_elements \* sizeof(int), cudaMemcpyHostToDevice);

// device to host

rc = cudaMemcpy(start, device\_mem, num\_elements \* sizeof(int), cudaMemcpyDeviceToHost);

// unified memory --> main() can work directly with device\_mem

cudaMallocManaged((void \*\*)&device\_mem, num\_elements \* sizeof(int));

// unified memory w/o malloc --> must know size at compile time

\_\_managed\_\_ int result[NUM\_ELEMENTS];

// copy from variable

cudaMemcpyFromSymbol(&host\_result, result, sizeof(start));

cudaMemcpyToSymbol(dResult, &result, sizeof(unsigned long long));

// how to get index using thread ID

// gridDim.xyz            blockDim.xyz

// blockIdx.xyz          threadIdx.xyz

// const unsigned long long id = threadIdx.x + blockIdx.x \* blockDim.x;

## OPEN\_MP

```
// variables declared before pragma are implicitly shared
// variables declared in the loop are implicitly private
pragma this and that

// parallel for loop with n_threads, reduce the variable `killed` using the `+` operation
#pragma omp parallel for num_threads(n_threads) reduction(+ : killed) collapse(2)

// explicit shared and private vars
#pragma omp parallel for shared(a, b, result) private(i, j, k)

// schedule --> how to break up loops
static distributes in order (block)
dynamic does round robin (cyclic/block cyclic) --> eg when different iteration might need different
and indeterminate amount of time
#pragma omp for schedule (static, chunksize)

// sections --> each `section` of code is executed in parallel
#pragma omp parallel {
    #pragma omp section{}
    #pragma omp section{}
}

// The nowait clause overrides any synchronization that would otherwise occur at the end of a
construct.

// only 1 thread executes this
any 1 thread --> #pragma omp single
master thread 0 --> #pragma omp master

// create a critical section
#pragma omp critical
{}

// barrier --> threads in a parallel region will not execute beyond the omp barrier until all other
threads in the team complete all explicit tasks in the region.
#pragma omp barrier

// atomic --> makes the next statement atomic
#pragma omp atomic
x[index[i]] += y;
```

## MPI

An intra-communicator refers to a single group, an inter-communicator refers to a pair of groups.

```
// send recv
int MPIAPI MPI_Send(
    _In_opt_ void *buf,
    int count,
    MPI_Datatype datatype,
    int dest,
    int tag,
    MPI_Comm comm
);

int MPIAPI MPI_Recv(
    _In_opt_ void *buf,
    int count,
    MPI_Datatype datatype,
    int source,
    int tag,
    MPI_Comm comm,
    _Out_ MPI_Status *status
);

// lsend lrecv
int MPIAPI MPI_Lsend(
    _In_opt_ void *buf,
    int count,
    MPI_Datatype datatype,
    int dest,
    int tag,
    MPI_Comm comm,
    _Out_ MPI_Request *request
);

int MPIAPI MPI_Lrecv(
    _In_opt_ void *buf,
    int count,
    MPI_Datatype datatype,
    int source,
    int tag,
    MPI_Comm comm,
    _Out_ MPI_Request *request
);
```

```

// scatter gather
int MPIAPI MPI_Scatter(
    _In_ void *sendbuf,
    int sendcount,
    MPI_Datatype sendtype,
    _Out_ void *recvbuf,
    int recvcnt,
    MPI_Datatype recvtpe,
    int root,
    MPI_Comm comm
);
int MPIAPI MPI_Gather(
    _In_ void *sendbuf,
    int sendcount,
    MPI_Datatype sendtype,
    _Out_opt_ void *recvbuf,
    int recvcnt,
    MPI_Datatype recvtpe,
    int root,
    MPI_Comm comm
);
int MPIAPI MPI_Allgather(
    _In_ void *sendbuf,
    _In_ int sendcount,
    _In_ MPI_Datatype sendtype,
    _Out_ void *recvbuf,
    int recvcnt,
    MPI_Datatype recvtpe,
    MPI_Comm comm
);

```

scatter sends `n` elements from the array in root to each receipient (by communicator)  
gather in main receives `n` elements from all members on communicator, puts into array  
allGather does gather + broadcast, all members have an array to write into

```

// broadcast --> uses a tree-like structure for efficiency
// root sends the data at &data. all others receive it at their &data
MPI_Bcast(
    void* data,
    int count,
    MPI_Datatype datatype,
    int root,
    MPI_Comm communicator)

// barrier
MPI_Barrier(MPI_Comm communicator)

```

```

// reduce --> reduction on a per-element basis, ie a[0] reduced and a[1] reduced separately
MPI_Reduce(
    void* send_data,
    void* recv_data, <-- only needed on root
    int count,
    MPI_Datatype datatype,
    MPI_Op op,
    int root,
    MPI_Comm communicator)
// MPI_Allreduce --> reduce followed by broadcast
MPI_Allreduce(
    void* send_data,
    void* recv_data,
    int count,
    MPI_Datatype datatype,
    MPI_Op op,
    MPI_Comm communicator)

```